

PROJECT 1: PREDICT THE HOUSING PRICES IN AMES (PART II)

You are asked to implement Lasso using Coordinate Descent on the Ames housing data.

Download the dataset, "Ames_data.csv", from the Resources page. The dataset has 2930 rows (i.e., houses) and 83 columns. Column 1 is "PID", the Parcel identification number, the last column is the response variable, "Sale_Price", and the remaining 81 columns are explanatory variables describing (almost) every aspect of residential homes.

The goal is to predict the final price of a home in Log scale.

Coordinate Descent for Lasso

- First write your own function or use the following function to solve the one-step Lasso for β_j , see page 37 of [lec_W3_VariableSelection.pdf].

```
one_step_lasso = function(r, x, lam){
  xx = sum(x^2)
  xr = sum(r*x)
  b = (abs(xr) - lam/2)/xx
  b = sign(xr)*ifelse(b>0, b, 0)
  return(b)
}
```

- Then write a function to implement CD for n_{iter} steps.

```
mylasso = function(X, y, lam, n.iter = 50, standardize = TRUE)
{
  # X: n-by-p design matrix without the intercept
  # y: n-by-1 response vector
  # lam: lambda value
  # n.iter: number of iterations
  # standardize: if True, center and scale X and y.

  # YOUR CODE
  # If standardize = TRUE, center and scale X and Y
  # record the corresponding means and sd

  # Initial values for residual and coefficient vector b
  b = rep(0, p)
  r = y

  for(step in 1:n.iter){
    for(j in 1:p){

      # YOUR CODE

      # 1) Update the residual vector
      # r <-- r + X[, j] * b[j]
      # r on the left: residual in blue on p37 of [lec_W3_VariableSelection.pdf]
      # r on the right: current residual

      # 2) Apply one_step_lasso to update beta_j
      # b[j] = one_step_lasso(r, X[, j], lam)

      # 3) Update the current residual vector
```

```

      # r <-- r - X[, j] * b[j]
    }
  }

  # YOUR CODE: scale back b and add intercept b0
  # For b0, check p13 of [lec_W3_VariableSelection.pdf].
  return(c(b0, b))
}

```

Note: in the script above, we run a fixed number of steps. `n.iter = 50` is just for illustration. You may need to set a bigger number. You can change it to a "while" loop to stop when some convergence criterion is satisfied.

- In practice, you need to add another loop outside to compute the Lasso solution for a range of lambda values.

How to pick the lambda sequence? First think about how to find the largest lambda value (denoted by `lambda_0`), i.e., when lambda is bigger than `lambda_0`, all coefficients are zero, and once lambda is smaller than `lambda_0`, one variable will enter the model. Then, set a sequence of lambda from `lambda_0` to a small number times `lambda_0`, equally spaced in log-scale.

In "mylasso" function above, the initial value of beta is always set to be all zero. You can try to use the fitted beta from a larger lambda value as the initial value for the next (smaller) lambda.

For this assignment, students are allowed to "cheat": you can use "glmnet" to figure out a proper range of lambda values, try them, and then fix a particular lambda value in your submission.

- Split the data into training and test with the `test.id` generated as follows. Fit your Lasso algorithm on the training data and pick the best lambda value by minimizing the RMSE of the logarithm of `Sale_Price` on the test data.

```
test.id = seq(1, 2930, by=3)
```

- You can choose to apply your Lasso algorithm only on a subset of features. For example, we removed the following variables when trying our algorithm: `Street`, `Utilities`, `Land_Slope`, `Condition_2`, `Roof_Matl`, `Heating`, `Pool_QC`, `Misc_Feature`, `Low_Qual_Fin_SF`, `Three_season_porch`, `Pool_Area`, `Misc_Val`, `Longitude` and `Latitude`.

What you need to submit?

Before the deadline ([Thursday, Oct 18, 11:30pm, Pacific Time](#)) please submit your R/Python code (.R or .py or zip).

How we evaluate your code?

Name your main file as **mymain_Lasso.R**. If you have multiple R files, upload the zip file. After unzipping your file, we will run the command `"source(mymain_Lasso.R)"` in a directory, in which there are only two files: `train.csv` and `test.csv`, where training and test are split using `test.id` described above.

- **train.csv**: subset the whole data "Ames_data.csv" with exactly the same 83 columns;
- **test.csv**: subset of whole data without the last column "Sale_Price", i.e., it has 82 columns.

After running your Rcode, we should see **ONE txt files** in the same directory named "mysubmission3.txt". **Submission File Format** and **Evaluation Procedure** are the same as Part I.

Your RMSE should be below 0.125.

The evaluation process for Python code is the same.
